

Predicting Economic Damage of Natural Disasters

Hyeonsu An, Max Niu

2026-04-20

Introduction

Project Motivation:

As global temperatures rise, climate-related disasters have increased in both frequency and severity over the past several decades ([World Wild Life](#)). Extreme weather events such as floods, droughts, hurricanes, etc. impose a huge human and real economic costs across the world. It is imperative that we understand the determinants of disaster severity (in terms of economic damages) for better policy planning and mitigative strategies.

The Emergency Events Database (EM-DAT), maintained by the Centre for Research on the Epidemiology of Disasters (CRED), keeps an internationally standardized data on natural disasters worldwide. The information includes disaster type, location, affected population, deaths, injuries, and economic damages.

Include discussion about work others have done related to your research question.

Research Question:

How do different disaster characteristics (such as geographic location, time of year, type, and total casualties incurred) predict the total economic damage caused by natural disasters in USD?

- Specifically, we examine whether disaster type, location, timing, duration, and measures of human impact (deaths, injuries, and affected population) are associated with higher reported economic losses.

Motivation:

We were first drawn to this research question because we were intrigued by this data; we were amazed by the granularity and the surprising breadth of information in this data. However, we realized that doing research on this matter could have real, practical uses.

First, identifying predictors of large economic damage is critical for both public and private institutions. Governments and international organizations allocate substantial resources towards mitigation and recovery efforts, yet the determinants of high-cost disasters vary across event types and regions. Same applies for insurance companies and hedge fund traders that try modeling key determinants relevant to a disaster for their comparative advantage against peers in the market.

By analyzing global disaster data, this project aims to quantify which observable characteristics are most strongly associated with economic loss. We hope that these results may provide insight into which types of disasters and impact measures are most economically destructive.

Hypothesis Reasoning:

Based on economic reasoning and relevant research, we expect that...

1. Disasters with higher number of affected individuals (absolute number, number of deaths and injuries) will have a positive relationship with economic damage.
2. Certain disaster types will incur systematically higher economic damage than that of others.
3. The level of damage dollars will differ based on countries because of their respective infrastructure that exists to mitigate damage.

Data Description:

This data was obtained from the Emergency Events Database, otherwise known as [EM-DAT](#), which is maintained by the Centre for Research on the Epidemiology of Disasters (CRED). The goal of the database is to compile standardized information on natural disasters that occur around the world. According to the information on metadata, an event is included in EM-DAT if one or more of the following criteria is met:

- 10+ people killed
- 100+ people affected
- Declaration of a state of emergency
- Solicitation of international assistance

The observations describe different disasters that occurred around the world. Of each disaster, data on the number of people injured, affected, homeless, or killed, the total economic damages, type of disaster, exact location, start date and end date of the disaster were collected.

The data dictionary can be found [here](#).


Limitations and Concerns

There are a considerable number of outliers in our data set with respect to total economic damage, duration in days, and many other variables. This motivated us to apply the logarithmic function to those variables. These outliers, however, should still be considered in the larger-picture analysis.

Exploratory Data Analysis

Response variable - log-transformed total economic damage in thousands of USD

Distribution of Economic Damage of Natural Disasters
from 2000 – 2025




min	1st_qu	median	mean	3rd_qu	max	sd
2.485	9.903	11.604	11.493	13.235	19.163	2.606

We will be working with log-transformed total economic damage in thousands of USD, because without the log transformation, the 123 outliers skew the data heavily to the right. The distribution of the logged response variable is unimodal and symmetric. The center is best represented by the mean, which is \$98,027,173.52 ($e^{11.493} * 1000$), and the spread is best represented with the standard deviation, which is \$13544.98 ($e^{2.606} * 1000$). There are outliers present in our distribution.

Predictors of interests-


1. disaster_type - the type of disaster (categorical)



disaster_type	median	IQR
Earthquake	11.951	3.283
Extreme temperature	12.546	0.000
Flood	11.379	3.429
Mass movement (wet)	9.903	3.229
Storm	11.546	3.459
Volcanic activity	11.542	0.592
Wildfire	12.559	1.746

Wildfires have the highest median log of total economic damage (in thousands of dollars), which is 12.559, and thus, normal total economic damage, which would be 284,604.84 thousands of dollars ($e^{12.559}$). The median of earthquakes is slightly lower, and then the median of volcanic activities, floods, and storms, which all have similar values, slightly below the median of earthquakes. Wet mass movements have the lowest median (log of) total economic damage. Storms, floods, and earthquakes have the highest variability in the (log of) total economic damage.


2. duration - duration in days of the disaster (quantitative)



```
# A tibble: 1 x 1
  r_squared
    <dbl>
1  0.00360
```


There seems to be a weak positive linear relationship between the log of duration in days with log of total economic damage (in thousands of dollars). The r-squared value is 0.003596655. There is a lot of variability in the response variable when the log of duration in days is low, but the variability decreases as the x value increases. There are outliers present throughout the graph.

Based off of our knowledge about how the duration of various disasters and how that should affect economic damage, we then decide to explore an interaction effect between `log_duration_days` and `disaster_type`.




After visualizing the potential interaction effect between log duration of days and type of disaster, it is evident that an increase in the log duration in days has a stronger effect on log total economic damage for storms and floods, compared to wildfires, volcanic activities, and wet mass movements, which is shown by the slightly steeper positive slope. We will consider putting this interaction effect into our model because of this visualization.

3. total_deaths - Total fatalities (deceased and missing combined) (quantitative)




There is a weak positive linear relationship between log of total deaths and economic damage with an r-squared value of 0.116. Disasters with higher fatalities generally correspond to higher economic losses (USD). However, the relationship is noisy and dispersed, indicating that deaths alone may not be a strong predictor of damage and that further analysis needs to be done.

4. no_injured - Number of people with physical injuries, trauma, or illness requiring immediate medical assistance due to the disaster (quantitative)



There is a weak positive linear relationship between log of total injuries and economic damage with an r-squared value of 0.18. This relationship is similar to `log_total_deaths`. Although slightly less, the relationship is noisy and dispersed, indicating that log number of injuries may not be a strong predictor of damage and that further analysis needs to be done.

5. `no_affected` - Number of people requiring immediate assistance due to the disaster (quantitative)



[1] 0.04448019

There is a weak positive linear relationship between log of total injuries and economic damage with an r-squared value of 0.044. This relationship is similar to `log_total_deaths` and `log_no_injured`, but the r-squared value is lower. This means that the relationship is more noisy and dispersed and that there is a stronger relationship between log total deaths and log number of injuries than log of number of people affected with log of total economic damages. We will keep this in mind while finding the best model.

Methodology

Determining the best model -

Based on our EDA, we were curious to see if a model with an interaction effect between `log_duration_days` and `disaster_type` would be a better model, so we conducted 10-fold cross validation to see which model performed better.

.metric	mean	n	std_err
rmse	2.240	10	0.044
rsq	0.227	10	0.037

.metric	mean	n	std_err
rmse	2.239	10	0.048
rsq	0.232	10	0.038

After conducting cross-validation, we observed similar values for rmse and r-squared for both models. Because the model without the interaction effect is the most parsimonious model, we are going to move forward with the model with log no_affected, log duration_days, disaster_type, log total_deaths, and log no_injured without any interaction effects.

term	estimate	std.error	statistic	p.value
(Intercept)	8.898	0.440	20.239	0.000
log(no_affected)	0.007	0.039	0.185	0.853
log_duration_days	0.285	0.107	2.667	0.008
disaster_typeExtreme temperature	-0.072	2.274	-0.032	0.975
disaster_typeFlood	-0.593	0.373	-1.590	0.112
disaster_typeMass movement (wet)	-1.089	0.673	-1.619	0.106
disaster_typeStorm	0.000	0.282	0.001	0.999
disaster_typeVolcanic activity	0.107	1.311	0.081	0.935
disaster_typeWildfire	1.179	0.620	1.902	0.058
log(total_deaths)	0.161	0.065	2.460	0.014
log(no_injured)	0.422	0.058	7.306	0.000

After displaying the model coefficients and p-values, we realized that the coefficient for log(no_affected) was not statistically significant at all because of its p-value of 0.853, which is much greater than 0.05. We decided to explore one last model, a model with all the previous predictors as additive effects but without log(no_affected). We evaluated the performance of these models through another round of 10-fold cross-validation, comparing RMSE and r-squared to determine the better model.

.metric	mean	n	std_err
rmse	2.240	10	0.044
rsq	0.227	10	0.037

.metric	mean	n	std_err
rmse	2.236	10	0.043
rsq	0.229	10	0.037

After evaluating the performance of the two models through cross-validation, we discovered that the values for RMSE and r-squared were similar again. Because the model without $\log(\text{no_affected})$ is more parsimonious and performed slightly better with a lower RMSE value of 2.236 compared to 2.240, we decided to choose the model without $\log(\text{no_affected})$ as our final model.

term	estimate	std.error	statistic	p.value
(Intercept)	8.952	0.325	27.506	0.000
log_duration_days	0.287	0.106	2.718	0.007
disaster_typeExtreme temperature	-0.037	2.264	-0.016	0.987
disaster_typeFlood	-0.596	0.372	-1.601	0.110
disaster_typeMass movement (wet)	-1.103	0.668	-1.651	0.099
disaster_typeStorm	0.004	0.281	0.013	0.990
disaster_typeVolcanic activity	0.106	1.310	0.081	0.936
disaster_typeWildfire	1.161	0.612	1.897	0.058
log(total_deaths)	0.166	0.059	2.787	0.005
log(no_injured)	0.423	0.057	7.365	0.000

Performance of final model on test data set:

```
# A tibble: 1 x 3
  .metric .estimator .estimate
  <chr>   <chr>         <dbl>
1 rmse    standard       2.59
```

Results

For the chosen model, the model diagnostics were all satisfied and there were no influential points in the data set. The final model has an RMSE of 2.592, meaning the predicted values of log total economic damage (in thousands of USD) deviate from the observed values by approximately 2.592, on average. **In other words, the predicted values for total economic damage deviate from the observed values by approximately \$13,356.** This is only slightly higher than the average cross-validated RMSE from the training set (2.236), which is expected when evaluating on unseen data, suggesting the model generalizes well and is appropriate for this data.

The statistically significant predictors in the model provide meaningful insight into what drives economic damage from natural disasters:

1. no_injured:

For each 1% increase in the number of people injured, the expected total economic damage increases by approximately 0.422%, holding all other variables constant, making it the strongest and most statistically significant predictor in the model.

2. no_deaths:

For each 1% increase in total fatalities, the expected total economic damage increases by approximately 0.161%, holding all other variables constant, providing further evidence that human casualties are positively associated with economic losses.

3. duration_days:

Disaster duration also plays a meaningful role: for each 1% increase in the duration of the disaster in days, the expected total economic damage increases by approximately 0.285%, holding all other variables constant, suggesting that longer-lasting disasters tend to cause progressively greater economic losses.

4. disaster_typeWildfire:

Wildfires are expected to multiply total economic damage by a factor of 3.193 ($e^{1.161} = 3.193$) compared to earthquakes (the baseline category), holding all other variables constant, though this result is only marginally significant ($p = 0.058$).

The coefficients from the model suggest that human impact measures, particularly the number of people injured and disaster duration, are the most important observable predictors of economic damage from natural disasters. The positive relationship between injuries and economic damage makes intuitive sense: a higher injury toll signals a more intense and widespread event, which simultaneously destroys physical infrastructure and strains public health systems — all of which translate into greater economic costs. The significance of disaster duration further supports this logic, as prolonged events allow for more cumulative destruction of assets and livelihoods.

Discussion

The most consistent finding from our project is that human impact measures (specifically the number of people injured and total fatalities) are the strongest predictors of economic damage. A 1% increase in injuries predicts approximately 0.422% higher economic losses, and a 1% increase in fatalities predicts approximately 0.161% higher losses, holding other variables constant. This aligns with our discussion in the introduction suggesting that casualties serve as a proxy for disaster intensity: more severe events damage infrastructure and harm public health systems, all of which result in greater economic losses. Disaster duration was also a meaningful predictor, with a 1% increase in duration predicting a 0.285% increase in expected damage, consistent with the intuition that prolonged events allow for greater cumulative destruction.

Notably, the number of people affected (no_affected) was not a useful predictor in our model ($p = 0.853$), which was somewhat unexpected given our initial hypothesis. We considered the possibility that its predictive value may also be largely captured by the injuries and fatalities variables already in the model, but this was disproven by a non-substantive VIF value ($x < 10$).

Regarding disaster type, wildfires were the significant level as a predictor ($p = 0.058$), with expected damages roughly 3.25 times higher than earthquakes (the baseline), holding all else constant. While storms, floods, and earthquakes are the most frequent disaster types in our data, they are not meaningfully different predictors of economic damage from one another after controlling for human casualties and duration.

Some limitations should be noted. First, the EM-DAT database is subject to reporting biases: disasters in wealthier, more documented countries may have more complete damage estimates, while events in lower-income regions may be systematically underreported, reducing the reliability of our response variable. Also, the economic damage figures themselves are estimates and may reflect insurance payouts rather than true total losses, reflective measurement error. Finally, we removed a number of ‘N/A’ data points; while we were mindful when cleaning the data, if the removed data points are systematically different in character than the rest of the data, its omission could render biases that we have not considered,


Given that several disaster type categories (such as volcanic activity, extreme temperature, and mass movement) had very few observations, future work could consider combining these smaller categories into a broader grouping to produce more reliable estimates and reduce model instability. Also, we could bring time into the model. Future analyses could also explore whether the predictive relationships identified here hold across different time periods, given the documented increase in climate-related disaster frequency and severity, which may mean that older observations are less informative for predicting damages from more recent disasters.

Appendix

EDA

Response Variable: log of total economic damage in thousands of USD

This is the distribution of the response variable without the log transformation. The distribution is unimodal and extremely right-skewed. The center is best represented by the median, which is \$109,576,000, and the spread is best represented with the IQR, which is \$539,750,000 (559,750,000 - 20,000,000). There are 123 (828-705) outliers present in the upper end of our distribution.



Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
12	20000	109576	1519107	559750	210000000

1. disaster_type - the type of disaster (categorical)


Storms, floods, and earthquakes are the most common disaster types, composing most of the data. Storms are the most common. Extreme temperature, wet mass movements, volcanic activity, and wildfires are the least common disaster types, with extreme temperature being the least common.

```
# A tibble: 7 x 1
  disaster_type
```

```

<chr>
1 Flood
2 Storm
3 Earthquake
4 Mass movement (wet)
5 Wildfire
6 Volcanic activity
7 Extreme temperature


```




2. duration - duration in days of the disaster (quantitative)

Similar to the response variable, all of the quantitative predictor variables in our data set had extreme right skew because of the abundance of outliers. This motivated us to log-transform all of them. The histogram of the log of the duration has right skew and is bimodal. This makes sense because a lot of the disasters only last 1 day long due to their inherent nature, such as earthquakes. The median is 1.099 log duration in days, and the IQR is 1.946 log duration in days. There are outliers in the upper end of our distribution.

Distribution of Days of Disasters



Distribution of Duration of Disasters




Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.000	0.000	1.099	1.187	1.946	5.817

3. total_deaths - Total fatalities (deceased and missing combined) (quantitative)

In the density plot above, we can see that the overall distribution of the log total fatalities is unimodal and has a distinct right skew. From the summary statistics, we can tell that while the median number of fatalities is 25 people, some grand disasters have a catastrophic number of deaths, with the highest number being 222570 people.


Statistic	Value
Min.	1.000
1st Qu.	6.000
Median	25.000
Mean	921.675
3rd Qu.	80.000
Max.	222570.000



4. no_injured - Number of people with physical injuries, trauma, or illness requiring immediate medical assistance due to the disaster (quantitative)

Statistic	Value
Min.	1.00
1st Qu.	14.00
Median	60.00
Mean	2338.36


Statistic	Value
3rd Qu.	266.25
Max.	366596.00



The story is similar to that of the previous predictor, `log_total_deaths`. There is a clear unimodal shape with a right skew, and we know that this skewness has been alleviated by the log transformation (so in reality, much larger). It makes sense that the median for `log_no_injured` is higher than the median for `log_total_deaths` at 60 people because logically there will always be more people injured than killed at any given natural disaster. This also applies to the maximum number being higher than that of `total_deaths` at 366,596 people.

5. `no_affected` - Number of people requiring immediate assistance due to the disaster (quantitative)


Statistic	Value
Min.	11
1st Qu.	7000
Median	60000
Mean	1325411
3rd Qu.	500250
Max.	150000000




In the density plot above, the distribution of the log-transformed values appears roughly unimodal, though there is still a right skew. Compared to fatalities, the distribution of affected individuals is more spread out, reflecting the wide range of disaster scales – from local events to those impacting entire countries. From the summary statistics, there is a gap between the median (60,000 people) and the mean (over 1.3 million people), indicating a big right tail driven by a small number of extremely large disasters. The maximum value of 150,000,000 people affected highlights the scale at which certain disasters can disrupt entire populations. This makes sense considering the nature of natural disasters: while many events are relatively small, large-scale disasters can impact vast populations through displacement and infrastructure damage.


Model Diagnostics

Relationship of Residuals and Fitted Values



Distribution of Residual Values





```
# A tibble: 0 x 11
# i 11 variables: log_total_damage_000_us <dbl>, log_duration_days <dbl>,
#   disaster_type <chr>, log(total_deaths) <dbl>, log(no_injured) <dbl>,
#   .fitted <dbl>, .resid <dbl>, .hat <dbl>, .sigma <dbl>, .cooksd <dbl>,
#   .std.resid <dbl>
```

- The linearity condition is satisfied because the points are randomly scattered in the plot of fitted values vs. residual values. We can not guess whether the residual would be positive or negative for any predicted value.
- The constant variance condition is satisfied because the vertical spread is approximately equal for each predicted value. Although the cluster of points appear to form a fan shape, there are points outside of the fan shape pattern that allow this condition to be satisfied.
- The normality condition is satisfied because the distribution is unimodal and roughly symmetric and thus, similar to a normal distribution.
- The independence condition is satisfied because we are told that a disaster is included in the data set if at least one of the following criteria are met: 10+ people killed, 100+ people affected, declaration of a state of emergency, and solicitation of international assistance. Each disaster is an independent event as well. Furthermore, in the graph displaying the distribution of the residuals for each disaster type, we see that there is no correlation between the residuals and disaster type.

To determine if there is multicollinearity among the variables, we calculated the variance inflation factor (VIF) for each predictor variable.

log_duration_days	disaster_typeExtreme	temperature
1.959074		1.037022
disaster_typeFlood	disaster_typeMass	movement (wet)
3.555234		1.237285
disaster_typeStorm	disaster_typeVolcanic	activity
2.453624		1.039175
disaster_typeWildfire		log(total_deaths)
1.398274		1.827523
log(no_injured)		
2.049733		

Because all of the VIF values are below 5, we concluded that multicollinearity is not prevalent among the variables.